

Qasim Tahir

Lahore, Pakistan | +92 3204045409 | qasimbtahir@gmail.com
github.com/Qasim-Tahir | linkedin.com/in/qasimtahir | qasimtahir.vercel.app

PROFESSIONAL SUMMARY

Aspiring AI Engineer with hands-on experience building production RAG pipelines, knowledge graphs, and LLM-powered automation systems. Focused on evaluation-driven development, retrieval architecture, and secure deployment of generative AI applications.

SKILLS

Languages: Python, JavaScript, TypeScript, SQL, HTML, CSS, Tailwind CSS

Technologies: FastAPI, Flask, TensorFlow, Keras, LlamaCpp, LangChain, LlamaIndex, LangGraph, ChromaDB, Qdrant, FAISS, Celery, Redis, Next.js, Node.js, Ragas, NER, Docling, uv, React, MongoDB, Docker, Azure, LLMs, Llama, HuggingFace, Lora, Qlora, Github, Linux, n8n, zapier

EDUCATION

BS ,Data Science, National University of Computer and Emerging Sciences

Sep 2021 – Present

- Data Structures and Algorithms, Database Systems, Natural Language Processing , Data Visualization and Analysis, Generative AI

PROJECTS

Developer, Automated Research Paper Summarizer & Presentation Generation

Jan 2026 – Present

- Engineered a 4-stage GraphRAG pipeline ingesting research papers into a Neo4j knowledge graph across 5 retrieval sources, scoring 82.46% composite on a custom 4-axis evaluation framework benchmarked against GPT-5.5 and Gemini 3.1 Pro via Ragas.
- Built a two-pass VLM figure enrichment module with 6-category classification, context-grounded prompts, and a 3-tier hallucination detection system with MD5 cross-document deduplication.
- Implemented a hybrid NER pipeline populating a Neo4j property graph with 6 entity types and 7 relationship types, enforcing quality via canonical deduplication, two-tier blacklisting, and regex rejection patterns.
- Developed a HyDE-augmented QA engine with 2-hop graph retrieval and BGE-reranker-v2-m3 cross-encoder reranking, served to a Next.js frontend featuring 3D force-directed graph visualization via Three.js.

Developer, AI Portfolio Engine (qasimtahir.vercel.app)

2026

- Engineered a production RAG chatbot using Qdrant and Groq over 10 project documentation files, delivering source-attributed answers with sub-second latency via a markdown-aware LangChain ingestion pipeline.
- Architected a 4-layer LLM security pipeline regex injection detection with unicode/homoglyph normalization, GPT-OSS-Safeguard-20b classification, persona-hardened prompts, and output sanitization, hardening a public-facing RAG endpoint against prompt injection.
- Deployed a serverless React frontend on Vercel with in-memory rate limiting, CSP and X-Frame-Options headers, and sub-4s response times within Vercel's 10s function constraint.

Developer, AI-Powered Invoice Processing & Automation System

2026

- Built an end-to-end n8n automation pipeline processing invoice emails via a FastAPI service and Vision-Language Model, extracting structured fields from unstructured attachments with conditional validation logic.
- Engineered a fault-tolerant error routing layer that isolates invalid or ambiguous invoices into a dedicated review flow, preventing silent failures in automated extraction.